

Investigating and Predicting COVID-19 Confirmed Cases in U.S. : Analysis and Conclusion


Team #5:

Kiran Brar

Olivia Alexander

Kimberly Segura

Number of Confirmed COVID-19 Cases in U.S.



Output Generation and Analysis

The background of the slide features a series of overlapping, wavy, organic shapes in various shades of orange and red. These shapes flow from the bottom left towards the top right, creating a sense of movement and depth. The colors range from a bright, sunny orange to a deep, vibrant red, with some areas appearing as lighter, more translucent washes of color.

Preprocessing Techniques

- **Power transformation:** BoxCox() is configurable data transform method that evaluate a suite of transforms automatically and select a best fit (optimizes lambda).
 - $\lambda = -1.0$ is a reciprocal transform.
 - $\lambda = -0.5$ is a reciprocal square root transform.
 - $\lambda = 0.0$ is a log transform.
 - $\lambda = 0.5$ is a square root transform.
 - $\lambda = 1.0$ is no transform.

```
8 us_Y = us_df["positive"]
9
10 us_box_total, lam = boxcox(us_Y)
11 print('Lambda: %f' % lam)
```

Lambda: 0.046134

Preprocessing Techniques

- **Standardization:** Transform data to mean of 0 and standard deviation of 1-- Gaussian transform.

```
#Standardize|
scaler = StandardScaler()

scaled_train_labels = scaler.fit_transform(train_labels.reshape(-1, 1))
scaled_test_labels = scaler.transform(test_labels.reshape(-1, 1))
```

- **Normalization:** Scale data to different/minimized range.


```
#Scale/Normalize Input
minmax_scaler = MinMaxScaler()

scaled_train_labels = minmax_scaler.fit_transform(train_labels.reshape(-1, 1))
scaled_test_labels = minmax_scaler.transform(test_labels.reshape(-1, 1))
```


Preprocessing Techniques

- **Differencing:** Removes trends and seasonality from a time series dataset.

```
raw_cases = us_df["positive"]  
  
first_level = difference(raw_cases)  
plt.plot(first_level);  
plt.title("Line plot of First-level differenced Covid 19 cases");
```




```
autocorrelation_plot(first_level)  
plt.title("Autocorrelation plot of First-level differenced Covid 19 cases");
```



```
third_level_diff = difference(second_level)  
autocorrelation_plot(third_level_diff)  
plt.title("Autocorrelation plot of third-level differencing on Covid 19 cases");  
print("Data is now stationary!")
```


Data is now stationary!



Third-level Differencing removed the autocorrelation from our dataset

Converting Time Series into Supervised Learning

Sliding Window / Lagged Values: Using a previous number of values (rather than time) to predict the following value



```
def create_dataset(X, y, time_steps=1):  
    Xs, ys = [], []  
    for i in range(len(X) - time_steps):  
        v = X.iloc[i:(i + time_steps)].values  
        Xs.append(v)  
        ys.append(y.iloc[i + time_steps])  
    return np.array(Xs), np.array(ys)  
  
timesteps = 7  
train_x, train_y = create_dataset(train_data, train_data, timesteps)  
test_x, test_y = create_dataset(test_data, test_data, timesteps)
```

Models:

The background of the slide features a series of overlapping, wavy bands in various shades of orange and red. The colors transition from a light, pale orange at the top to a deep, vibrant red at the bottom, creating a sense of depth and movement. The waves are smooth and fluid, resembling a stylized landscape or a cross-section of a layered material.

Linear Regression

- Used sliding window approach
- **Best Linear Regression model:**


Lagged Values: 6 days

Box Cox: Yes

Differencing: No

Scaling/Normalizing: MinMaxScaler()

Linear Regression: Prediction Plots



MSE	0.0000402
SMAPE	0.84284873

Linear Regression Model: Different Lagged Values

Linear Regression with Different Preprocessing Techniques										
	Model 0	Model 1.1	Model 1.2	Model 1.3	Model 1.4	Model 1.5	Model 1.6	Model 1.7	Model 1.8	Model 2
Preprocessing										
Lagged Value	None	1	2	3	4	5	6	7	8	6
Box Cox /Log	None	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes
Scaling/Normalizing	None	Minmax()	Minmax()	Minmax()	Minmax()	Minmax()	Minmax()	Minmax()	Minmax()	Minmax(-1,1)
MSE (Test)	49347904776	0.001348057	0.000220618	0.000107409	8.45E-05	5.36E-05	4.02E-05	6.44E-05	6.61E-05	0.000160831
SMAPE (Test)	120.5941151	16.37057734	3.737763698	2.002087713	1.512598561	1.055856231	0.84284873	0.987350005	0.97504241	7.03589239
Notes:	LR on raw set						lowest error	error increasing again		auto minmax is better

Linear SVR

- Used sliding window approach
- **Best Linear SVR model:**


Lagged Values: 7 days

Box Cox: Yes

Differencing: No

Scaling/Normalizing: MinMaxScaler()

Linear SVR: Prediction Plots




MSE	0.00009108727842
SMAPE	1.145
RMSE	0.009544

Nonlinear SVR with RBF Kernel

Nonlinear SVR model (with lowest MSE):

- Lagged Values: 7 days
- Box Cox: Yes
- Differencing: No
- Scaling/Normalizing: MinMaxScaler()

Nonlinear SVR: Prediction Plots



MSE	0.000135
SMAPE	1.455
RMSE	0.011619

Linear SVR: Effect of Lag Value

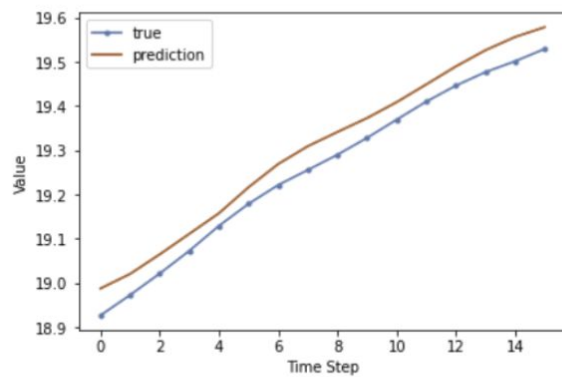
	Linear SVR: Performance for Different Lag Values			
	Model 3.1	Model 3.2	Model 3.3	Model 3.4
Features:				
Preprocessing				
Lagged Value	3	5	10	8
Box Cox /Log	Yes	Yes	Yes	Yes
Differencing	No	No	No	No
Scaling/Normaliz	Minmax(0,1)	Minmax(0,1)	Minmax(0,1)	Minmax(0,1)
Model				
Objective	Linear	Linear	Linear	Linear
Kernel	-	-	-	-
Hyperparamters				
C	100	1000	100	1
Epsilon	0.0005	0.0001	0.0001	0.001
Gamma	-	-	-	-
MSE	0.001026	0.000804	0.0001094	0.000114
SMAPE	5.068	4.284	1.072	1.2952
RMSE	0.0320312	0.0283549	0.0104594	0.0106771

Linear SVR: Effect of Scaling vs. Standardizing


	LinearSVR Model Performances		
	Model 5	Model 5.1	Model 5.5
Features:			
Preprocessing			
Lagged Value	7	7	7
Box Cox /Log	No	No	No
Differencing	No	No	No
Scaling/Normali	Minmax(0,1)	Minmax(-1,1)	StandardScaler
Hyperparamters			
C	1000	100	1000
Epsilon	0.001	0.0001	0.0001
MSE	0.000276	0.00049046	0.0033577
SMAPE	2.419158	9.68705	11.82743
RMSE	0.0166132	0.0221463	0.0579457

	Support Vector Regression: Performance Results					
	Model 1	Model 2	Model 5.5	Model 6	Model 7	Model 8
Features:						
Preprocessing						
Lagged Value	7	7	7	7	7	7
Box Cox /Log	Yes	Yes	No	Yes	Yes	No
Differencing	No	No	No	No	No	No
Scaling/Normaliz	Minmax(0,1)	Minmax(0,1)	StandardScaler	None	None	None
Model						
Objective	Linear	Non-linear	Linear	Linear	Nonlinear	Linear
Kernel	-	RBF	-	-	rbf	-
Hyperparamters						
C	1000	1000	1000	1	1000	1
Epsilon	0.0001	0.0001	0.0001	0.0005	0.001	0.5
Gamma	-	0.005	-		0.001	
MSE	0.000091087271	0.000135	0.0033577	0.0246214	0.00213	307124112.3
SMAPE	1.145	1.455	11.82743	0.810457	0.236	1.256246775
RMSE	0.009544	0.011619	0.0579457	0.1569121	0.0461519	17524.95684


Model 5.5



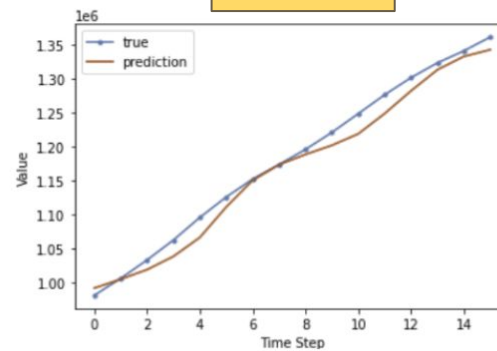
Model 6



Model 7




Model 8



XGBoost :

Approach 1 (Use month, day as features)



	XGBoost: Using Date Features		
	Model 0	Model 0.5	Model 1
Features:	Month, day	Month, day	Month, day
Preprocessing			
Box Cox /Log	No	Yes	Yes
Scaling/Normali	Minmax	Minmax	StandardScaler
MSE (Test)	883237648.7	168	2.199
SMAPE (Test)	69	21.56	6.789
RMSE(Test)	29719.3144	12.96148	1.4829

XGBoost :

Approach 2 (Lagged Values as Input)

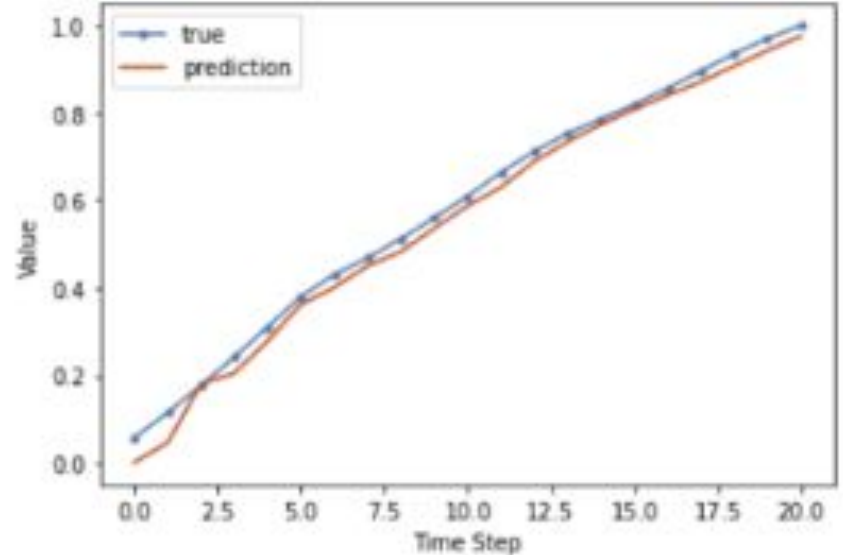
- Used sliding window approach
- Best model

Lagged Values: 1

Box Cox: Yes

Differencing: No

Scaling/Normalizing: MinMaxScaler()




XGBoost : Effect of Lagged Values

- From our experiments, we found that the smaller lag values increased performance for XGBoost.


	XGBoost Model Performances with Lagged Values Features and No Differencing							
	Model 2	Model 2.1	Model 2.3*	Model 2.5	Model. 2.6	Model. 2.7	Model 2.8	Model 2.9
Features:	LagVal	LagVal	LagVal	LagVal	LagVal	LagVal	LagVal	LagVal
Preprocessing								
Lagged Value	7	7	7	10	3	5	2	1
Box Cox /Log (0	Yes	Yes	Yes	Yes	Yes	yes	Yes	Yes
Scaling/Normali	Minmax [0,1]	Minmax[-1,1]	Minmax [0,1]	Minmax [0,1]	Minmax [0,1]	Minmax [0,1]	Minmax [0,1]	Minmax [0,1]
MSE (Test)	0.02119	0.022	0.01072	0.04252	0.0049	0.011	0.002751	0.0009066
SMAPE (Test)	22	23.53	17.649	30.882	12.370	18.46	13.54425	17.22109
RMSE	0.1455679	0.148324	0.1035374	0.2062038	0.07	0.1048809	0.05245	0.0301098

XGBoost: Effect of Differencing


Model 3




Model 3.1



Model 3.5



Model 3.8




XGBoost: Effect of Differencing

	XGBoost Models using Differencing Preprocessing Techniques						
	Model 3	Model 3.1	Model 3.5	Model 3.6	Model 3.7	Model 3.8	Model 3.9
Features:	LagVal	LagVal	LagVal	LagVal	LagVal	LagVal	LagVal
Preprocessing							
Lagged Value	3	3	3	3	1	1	1
Box Cox /Log	None	None	Yes	Yes	Yes	Yes	Yes
Differencing Level	3	1	1	2	2	3	3
Scaling/Normalization	None	None	Minmax [0,1]	Minmax [0,1]	Minmax [0,1]	Minmax [0,1]	StandardScaler
MSE (Test)	192.291	200	68.064	132.445	154.858	46.12	174.962
SMAPE (Test)	1.6506	0.08078	0.0883	0.35	0.425	0.082	2.047
RMSE	13.8669168	14.1421356	8.2500909	11.5084751	12.4441954	6.7911707	13.2273202

LSTM Best Model

```
] model = keras.Sequential()  
model.add(keras.layers.LSTM(50, input_shape=(X_train.shape[1], X_train.shape[2]), return_sequences = True))  
model.add(keras.layers.Dropout(0.1))  
model.add(keras.layers.LSTM(units = 50, return_sequences = True))  
model.add(keras.layers.Dropout(0.1))  
model.add(keras.layers.Dense(1))  
model.compile(loss='mean_squared_error', optimizer = keras.optimizers.Adam(0.1))
```


```
#Diagnostic Line Plots for Loss and Validation loss of the data  
>lt.plot(history.history['loss'], label='train')  
>lt.plot(history.history['val_loss'], label='test')  
>lt.legend();
```



□

Evaluation on test data: MSE
1/1 [=====] - 0s 970us/step - loss: 0.0082
0.0081513412296772

Evaluation on test data: SMAPE
117.20786382367955



LSTM: Effect of Lag Value and Differencing

	LSTM Performance of Different Lag Values (No Differencing)				LSTM Performance of Different Lag Values (with Differencing)			
Trial	1	2	3	4	5	6	7	8
Features:								
Preprocessing								
Lagged Value	3	3	5	7	3	5	7	10
Box Cox /Log	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes
Differencing	No	No	No	No	Yes	Yes	Yes	Yes
Scaling/Normali	MinMax(-1,1)	MinMax(0,1)	MinMax(0,1)	MinMax(0,1)	MinMax(0,1)	MinMax(0,1)	MinMax(0,1)	MinMax(0,1)
MSE (Test)	0.161956	0.043926	0.023220	0.044450	0.030001	0.010664	0.009857	0.008716
SMAPE (Test)	917.383506	448.378769	283.922014	359.388216	363.644981	187.775853	159.572204	121.370556
RMSE	0.026230	0.001929	0.000539	0.001976	0.000900	0.000114	0.000097	0.000076

Compare output against Hypothesis

Hypothesis: We predicted long short term networks (LSTM) to yield the best accuracy, followed by XGBoost, SVR, and linear regression.

Output:



- Linear Regression and SVR had overall best performance(with preprocessing techniques).
- SVR had better performance than XGBoost and LSTM on **raw** data (in general)
- LSTM had the worst predictions

Abnormal Case Explanation

- Differencing did not always improve performance (only slightly helped for LSTM)
- SVR had surprisingly accurate predictions on untransformed data, in comparison to other models.
- Smaller lag values improved performance for XGBoost, whereas larger lag values improved performance for LR, SVR, and LSTM
- Minmax Scaler better performance than StandardScaler

Discussion

- We found lagged values, and power transformations (Box Cox) to be incredibly helpful to improve model performance.
- We also found LSTM had poor and unstable performance. This can be due to the small dataset or lack of time to tune the number of layers/neurons
 - Same model yielding different results:



Conclusion/ Recommendation

The background of the slide features a series of overlapping, wavy, organic shapes in various shades of orange and red. These shapes flow from the bottom left towards the top right, creating a sense of movement and depth. The colors range from a bright, sunny orange to a deep, vibrant red, with some areas appearing as lighter, more translucent washes over darker tones.

Summary/Conclusion

- Don't underestimate the importance of preprocessing or feature engineering!
- Important to test simple models first!
- For deep learning, you need a lot of data and it's very time consuming to pick and choose different number of layers and neurons.

Recommendations for Future Studies

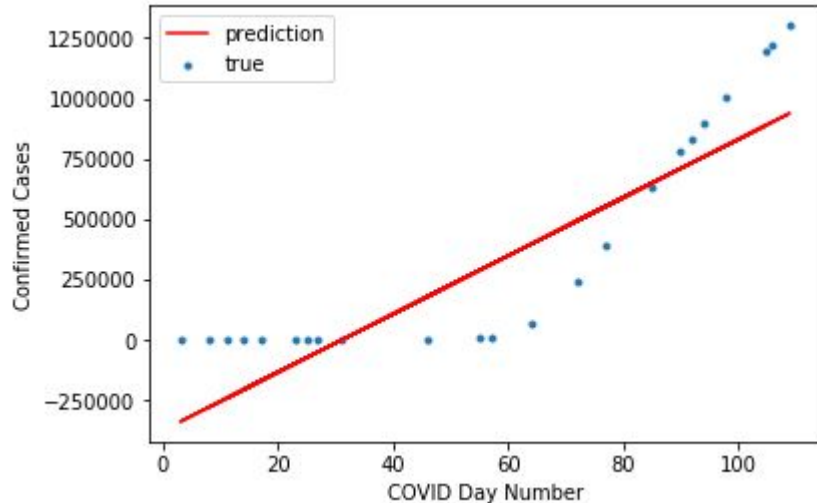
- Our work can be furthered by predicting number of deaths or number of recovered.
- Our models can be extended to take into consideration the health capacity as well as social restrictions for each in order to make better future forecasting.

Demonstration:

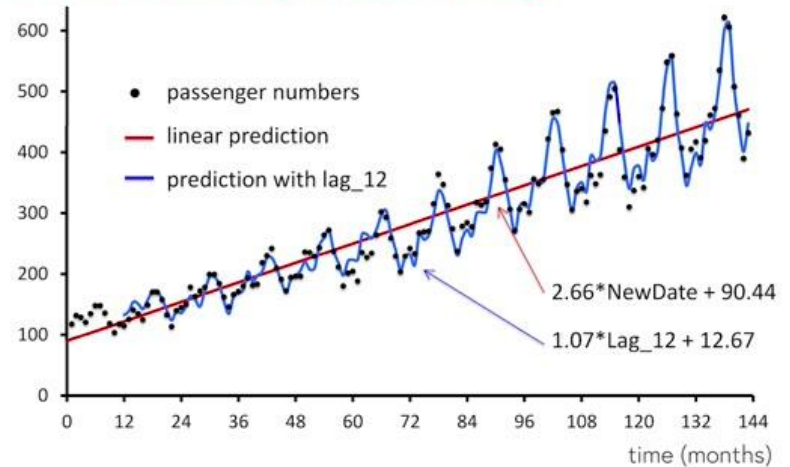
- Model Training and Evaluation
 - EDA
- 

Linear Regression Model: Nonlinear Prediction

Linear Prediction vs Nonlinear Prediction: Transform predictor X from time to sliding window. Nonlinear functional form but model still linear in parameters.



Time series: Linear regression with lags



Extra Reference Slides - LSTM approach 1

MODEL

```
[ ] #  
    lstm_model = Sequential()  
    lstm_model.add(LSTM(units = 50, return_sequences = True, input_shape = (n_input, n_features)))  
    lstm_model.add(Dropout(0.2))  
    lstm_model.add(LSTM(units = 50, return_sequences = True))  
    lstm_model.add(Dropout(0.2))  
    lstm_model.add(LSTM(units = 50))  
    lstm_model.add(Dropout(0.2))  
    lstm_model.add(Dense(units = 1))  
    lstm_model.compile(optimizer = 'adam', loss = 'mean_squared_error')  
    history = lstm_model.fit(generator, epochs = 100)
```

Epoch 14/100

Model: "sequential_1"

Layer (type)	Output Shape	Param #
=====	=====	=====
lstm_1 (LSTM)	(None, 10, 50)	10400
dropout_1 (Dropout)	(None, 10, 50)	0
lstm_2 (LSTM)	(None, 10, 50)	20200
dropout_2 (Dropout)	(None, 10, 50)	0
lstm_3 (LSTM)	(None, 50)	20200
dropout_3 (Dropout)	(None, 50)	0
dense_1 (Dense)	(None, 1)	51
=====	=====	=====

Total params: 50,851

Trainable params: 50,851

Non-trainable params: 0

Extra Reference Slides - LSTM approach 1

TimeseriesGenerator

```
] #LAG preprocessing to frame a sequence as a supervised learning problem
#returns a sequence of overlapping windows
#batch_size = # of samples to return on each iteration
from tensorflow.keras.preprocessing.sequence import TimeseriesGenerator

n_input = 10 # lag
n_features = 1
generator = TimeseriesGenerator(scaled_train_data, scaled_train_data, length = n_input, batch_size = 1)
for i in range(len(generator)):
    x, y = generator[i]
    #print('%s => %s' % (x, y))
```